



</>

Rlm integration spec · MD

Copy



Context Foundry × Recursive Language Models (RLM) Integration Spec

Purpose: Enhance Context Foundry's query/reasoning pipeline with RLM principles

Status: Draft for architectural review **Date:** January 2026

Executive Summary

This spec proposes integrating Recursive Language Model (RLM) principles into Context Foundry's existing tri-memory architecture. The core insight from the RLM paper (Zhang et al., MIT CSAIL, Dec 2025): **treat long contexts as part of an external environment** that LLMs can programmatically interact with, rather than stuffing everything into the context window.

Context Foundry already implements a version of this philosophy with its tri-memory separation. RLM enhancement would allow cognitive agents to **write code to navigate memory** and **recursively sub-query for verification** — dramatically improving performance on complex multi-hop queries across ADQ's 50+ portfolio companies.

Model Strategy: The architecture is model-agnostic, supporting both Anthropic (Claude 4.5 family) and OpenAI (GPT-5 family). Recommended default: Claude Sonnet 4.5 for root orchestration + Claude Haiku 4.5 for sub-queries, maintaining consistency with CF's existing use of Claude for vision tasks.

Current Architecture (Baseline)

User Query



QueryClassifier

Classify intent, check data sufficiency





EntityResolver

3-stage: exact → semantic → fuzzy

